

new_project_visa

libraries

```
library(randomForest)
```

```
## randomForest 4.6-12
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
library(party)
```

```
## Loading required package: grid
```

```
## Loading required package: mvtnorm
```

```
## Loading required package: modeltools
```

```
## Loading required package: stats4
```

```
## Loading required package: strucchange
```

```
## Loading required package: zoo
```

```
##  
## Attaching package: 'zoo'
```

```
## The following objects are masked from 'package:base':  
##  
## as.Date, as.Date.numeric
```

```
## Loading required package: sandwich
```

```
library(dplyr)
```

```
##  
## Attaching package: 'dplyr'
```

```
## The following object is masked from 'package:randomForest':  
##  
## combine
```

```
## The following objects are masked from 'package:stats':  
##  
## filter, lag
```

```
## The following objects are masked from 'package:base':  
##  
## intersect, setdiff, setequal, union
```

```
library(tidyr)  
library(sqldf)
```

```
## Loading required package: gsubfn
```

```
## Loading required package: proto
```

```
## Loading required package: RSQLite
```

```
library(caret)
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```
##  
## Attaching package: 'ggplot2'
```

```
## The following object is masked from 'package:randomForest':  
##  
## margin
```

```
library(glmnet)
```

```
## Loading required package: Matrix
```

```
##  
## Attaching package: 'Matrix'
```

```
## The following object is masked from 'package:tidyr':  
##  
## expand
```

```
## Loading required package: foreach
```

```
## Loaded glmnet 2.0-13
```

```
library(car)
```

```
## Warning: package 'car' was built under R version 3.4.3
```

```
##  
## Attaching package: 'car'
```

```
## The following object is masked from 'package:dplyr':  
##  
##      recode
```

```
library(ROCR)
```

```
## Loading required package: gplots
```

```
##  
## Attaching package: 'gplots'
```

```
## The following object is masked from 'package:stats':  
##  
##      lowess
```

```
library(e1071)  
library(gbm)
```

```
## Loading required package: survival
```

```
##  
## Attaching package: 'survival'
```

```
## The following object is masked from 'package:caret':  
##  
##      cluster
```

```
## Loading required package: splines
```

```
## Loading required package: parallel
```

```
## Loaded gbm 2.1.3
```

```
library(class)  
library(MASS)
```

```
##  
## Attaching package: 'MASS'
```

```
## The following object is masked from 'package:dplyr':  
##  
##      select
```

```
library(readr)  
library(tree)  
library(readr)  
h1b_kaggle <- read_csv("/Users/Pablo/Documents/IIT/Machine Learning/Project/h1b_kaggle.csv")
```

```
## Warning: Missing column names filled in: 'X1' [1]
```

```
## Parsed with column specification:  
## cols(  
##   X1 = col_integer(),  
##   CASE_STATUS = col_character(),  
##   EMPLOYER_NAME = col_character(),  
##   SOC_NAME = col_character(),  
##   JOB_TITLE = col_character(),  
##   FULL_TIME_POSITION = col_character(),  
##   PREVAILING_WAGE = col_double(),  
##   YEAR = col_integer(),  
##   WORKSITE = col_character(),  
##   lon = col_double(),  
##   lat = col_double()  
## )
```

Preprocessing data

```
#find rows where conditions are true where CASE_STATUS is CERTIFIED, CERTIFIED-WITHDRAWN that is treated as CERTIFIED or DENIED in the year 2014, 2015 and 2016
```

```
myData = filter(hlb_kaggle, hlb_kaggle$CASE_STATUS %in% c('CERTIFIED','DENIED', 'CERTIFIED-WITHDRAWN') & (hlb_kaggle$YEAR == 2016 | hlb_kaggle$YEAR == 2015 | hlb_kaggle$YEAR == 2014))
```

```
#Keep only complete cases
```

```
myData = myData[complete.cases(myData),]
```

```
hlbData = myData
```

```
#Eliminate columns case#, employer, job title, long, lat
```

```
hlbData[,c(1,3,5,10,11)]=NULL
```

```
#Create a new column called worksite to keep only state
```

```
hlbData=separate(data = hlbData, col = WORKSITE, into = c("CITY", "STATE"), sep = ",")
```

```
#Create a new column to save occupations
```

```
hlbData$occ=NA
```

```
#Keep occupations containing the keyword and set the new occupation
```

```
hlbData$occ[grepl("engineer",hlbData$SOC_NAME, ignore.case = T)]= "ENGINEER"
```

```
hlbData$occ[grepl("manager",hlbData$SOC_NAME, ignore.case = T)]= "MANAGER"
```

```
hlbData$occ[grepl("technician",hlbData$SOC_NAME, ignore.case = T)]= "TECHNICIAN"
```

```
hlbData$occ[grepl("teacher",hlbData$SOC_NAME, ignore.case = T)]= "TEACHER"
```

```
hlbData$occ[grepl("executive",hlbData$SOC_NAME, ignore.case = T)]= "EXECUTIVE"
```

```
hlbData$occ[grepl("accountant",hlbData$SOC_NAME, ignore.case = T)]= "ACCOUNTANT"
```

```
hlbData$occ[grepl("actor",hlbData$SOC_NAME, ignore.case = T)]= "ACTOR"
```

```
hlbData$occ[grepl("advertising",hlbData$SOC_NAME, ignore.case = T)]= "ADVERTISING"
```

```
hlbData$occ[grepl("lawyer",hlbData$SOC_NAME, ignore.case = T)]= "LAWYER"
```

```
hlbData$occ[grepl("financial",hlbData$SOC_NAME, ignore.case = T)]= "FINANCIAL"
```

```
hlbData$occ[grepl("arquitect",hlbData$SOC_NAME, ignore.case = T)]= "ARQUITECT"
```

```
hlbData$occ[grepl("programmer",hlbData$SOC_NAME, ignore.case = T)]= "SOFTWARE"
```

```
hlbData$occ[grepl("software",hlbData$SOC_NAME, ignore.case = T)]= "SOFTWARE"
```

```
hlbData$occ[grepl("computer",hlbData$SOC_NAME, ignore.case = T)]= "SOFTWARE"
```

```
hlbData$occ[grepl("developer",hlbData$SOC_NAME, ignore.case = T)]= "SOFTWARE"
```

```
hlbData$occ[grepl("analyst",hlbData$SOC_NAME, ignore.case = T)]= "ANALYST"
```

```
hlbData$occ[grepl("scien",hlbData$SOC_NAME, ignore.case = T)]= "SCIENTIST"
```

```
hlbData$occ[grepl("specialist",hlbData$SOC_NAME, ignore.case = T)]= "SPECIALIST"
```

```
hlbData$occ[grepl("animal",hlbData$SOC_NAME, ignore.case = T)]= "ANIMAL RELATED"
```

```
hlbData$occ[grepl("athlet",hlbData$SOC_NAME, ignore.case = T)]= "ATHLETE"
```

```
hlbData$occ[grepl("cook",hlbData$SOC_NAME, ignore.case = T)]= "COOK"
```

```
hlbData$occ[grepl("chef",hlbData$SOC_NAME, ignore.case = T)]= "COOK"
```

```
hlbData$occ[grepl("admin",hlbData$SOC_NAME, ignore.case = T)]= "ADMINISTRATIVE"
```

```
#Eliminate columns SOC_NAME and CITY
```

```
hlbData$SOC_NAME=NULL
```

```
hlbData$CITY= NULL
```

```
#Removing states with low count
```

```
a=sqlDF("select count(*) cc, STATE from 'hlbData' group by STATE")
```

```
b=sqlDF("select * from a where cc>2000 AND STATE <> ' NA'")
```

```
hlbData$STATE=ifelse(hlbData$STATE %in% b$STATE,hlbData$STATE,NA)
```

```

#Convert the dependent variable to binary
h1bData$CASE_STATUS=ifelse(h1bData$CASE_STATUS %in% c("CERTIFIED-WITHDRAWN", "CERTIFIED"), "1", "0")

#Converting categorical variables into factors
h1bData[,c(-3)]= lapply(h1bData[,c(-3)], as.factor)
h1bData = h1bData[complete.cases(h1bData),]

#Using years 2014 and 2015 as training and 2016 as test
data.test = h1bData[0:557090, ]
data.train = h1bData[557091:1522982,]
data.test = data.test[, -4]
data.train = data.train[, -4]

```

LOGISTIC REGRESSION USING ALL PREDICTORS

```

#Fitting the model on the training dataset
h1bglm.train.fit =glm(CASE_STATUS~., family=binomial(link = logit), data = data.train
)

```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```

#Finding Prdicitons on Testing set
prediction=predict(h1bglm.train.fit, newdata=data.test, type="response")
View(prediction)

#Threshold
pred_class = vector()
pred_class[prediction<0.5]=0
pred_class[prediction>=0.5]=1

#confusion matrix
confusionMatrix(pred_class, data.test$CASE_STATUS,positive="1")


```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction      0      1
##           0      51      1
##           1    6650 550388
##
##           Accuracy : 0.9881
##           95% CI : (0.9878, 0.9883)
##           No Information Rate : 0.988
##           P-Value [Acc > NIR] : 0.2719
##
##           Kappa : 0.0149
##           Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.999998
##           Specificity : 0.007611
##           Pos Pred Value : 0.988062
##           Neg Pred Value : 0.980769
##           Prevalence : 0.987971
##           Detection Rate : 0.987970
##           Detection Prevalence : 0.999907
##           Balanced Accuracy : 0.503804
##
##           'Positive' Class : 1
##
```

```
#levels(data.test$CASE_STATUS)
```

```
#ROC Curve
```

```
pred = prediction(pred_class, data.test$CASE_STATUS)
perf = performance(pred,"tpr","fpr")
plot(perf)
```



```
#Area Under the Curve  
auc.tmp = performance(pred,"auc");  
auc = as.numeric(auc.tmp@y.values)  
auc
```

```
## [1] 0.5038045
```

LOGISTIC REGRESSION USING ONLY FULL_TIME_POSITION AND PREVAILING_WAGE

```
#Fitting the model on the training dataset  
h1bglm.train.fit =glm(CASE_STATUS~.-STATE-occ, family=binomial(link = logit), data =  
data.train)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```



```

#Finding Prdicitons on Testing set
prediction=predict(h1bglm.train.fit, newdata=data.test, type="response")

#Threshold
pred_class = vector()
pred_class[prediction<0.5]=0
pred_class[prediction>=0.5]=1

#confusion matrix
confusionMatrix(pred_class, data.test$CASE_STATUS,positive="1")

```

```


## Confusion Matrix and Statistics
##
##           Reference
## Prediction      0      1
##           0      51      1
##           1    6650 550388
##
##           Accuracy : 0.9881
##           95% CI : (0.9878, 0.9883)
##      No Information Rate : 0.988
##      P-Value [Acc > NIR] : 0.2719
##
##           Kappa : 0.0149
##  McNemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.999998
##           Specificity : 0.007611
##           Pos Pred Value : 0.988062
##           Neg Pred Value : 0.980769
##           Prevalence : 0.987971
##           Detection Rate : 0.987970
##      Detection Prevalence : 0.999907
##           Balanced Accuracy : 0.503804
##
##           'Positive' Class : 1
##

```

```

#ROC Curve
pred = prediction(pred_class, data.test$CASE_STATUS)
perf = performance(pred,"tpr","fpr")
plot(perf)

```



```
#Area Under the Curve  
auc.tmp = performance(pred,"auc");  
auc = as.numeric(auc.tmp@y.values)  
auc
```

```
## [1] 0.5038045
```

LOGISTIC REGRESSION USING ONLY FULL_TIME_POSITION, PREVAILING_WAGE AND STATE

```
#Fitting the model on the training dataset  
h1bglm.train.fit =glm(CASE_STATUS~.-occ, family=binomial(link = logit), data = data.t  
rain)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```

#Finding Prdicitons on Testing set
prediction=predict(h1bglm.train.fit, newdata=data.test, type="response")

#Threshold
pred_class = vector()
pred_class[prediction<0.5]=0
pred_class[prediction>=0.5]=1

#confusion matrix
confusionMatrix(pred_class, data.test$CASE_STATUS,positive="1")

```

```


## Confusion Matrix and Statistics
##
##           Reference
## Prediction      0      1
##           0      51      1
##           1    6650 550388
##
##           Accuracy : 0.9881
##           95% CI : (0.9878, 0.9883)
##      No Information Rate : 0.988
##      P-Value [Acc > NIR] : 0.2719
##
##           Kappa : 0.0149
##  McNemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.999998
##           Specificity : 0.007611
##           Pos Pred Value : 0.988062
##           Neg Pred Value : 0.980769
##           Prevalence : 0.987971
##           Detection Rate : 0.987970
##      Detection Prevalence : 0.999907
##           Balanced Accuracy : 0.503804
##
##           'Positive' Class : 1
##

```

```

#ROC Curve
pred = prediction(pred_class, data.test$CASE_STATUS)
perf = performance(pred,"tpr","fpr")
plot(perf)

```



```
#Area Under the Curve  
auc.tmp = performance(pred,"auc");  
auc = as.numeric(auc.tmp@y.values)  
auc
```

```
## [1] 0.5038045
```

LOGISTIC REGRESSION USING ONLY FULL_TIME_POSITION, PREVAILING_WAGE AND OCCUPATION

```
#Fitting the model on the training dataset  
h1bglm.train.fit =glm(CASE_STATUS~.-STATE, family=binomial(link = logit), data = dat  
a.train)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```

#Finding Prdicitons on Testing set
prediction=predict(h1bglm.train.fit, newdata=data.test, type="response")

#Threshold
pred_class = vector()
pred_class[prediction<0.5]=0
pred_class[prediction>=0.5]=1

#confusion matrix
confusionMatrix(pred_class, data.test$CASE_STATUS,positive="1")

```

```


## Confusion Matrix and Statistics
##
##           Reference
## Prediction      0      1
##           0      51      1
##           1    6650 550388
##
##           Accuracy : 0.9881
##           95% CI : (0.9878, 0.9883)
##      No Information Rate : 0.988
##      P-Value [Acc > NIR] : 0.2719
##
##           Kappa : 0.0149
##  Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.999998
##           Specificity : 0.007611
##           Pos Pred Value : 0.988062
##           Neg Pred Value : 0.980769
##           Prevalence : 0.987971
##           Detection Rate : 0.987970
##      Detection Prevalence : 0.999907
##           Balanced Accuracy : 0.503804
##
##           'Positive' Class : 1
##

```

```

#ROC Curve
pred = prediction(pred_class, data.test$CASE_STATUS)
perf = performance(pred,"tpr","fpr")
plot(perf)

```



```
#Area Under the Curve  
auc.tmp = performance(pred,"auc");  
auc = as.numeric(auc.tmp@y.values)  
auc
```

```
## [1] 0.5038045
```

LDA USING USING ALL PREDICTORS

```
lda.fit=lda(CASE_STATUS~., data=data.train)  
lda.predict = predict(lda.fit, data.test, type = "prob")  
confusionMatrix(lda.predict$class, data.test$CASE_STATUS,positive="1")
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction      0      1
##           0      105      360
##           1     6596    550029
##
##           Accuracy : 0.9875
##           95% CI : (0.9872, 0.9878)
##           No Information Rate : 0.988
##           P-Value [Acc > NIR] : 0.9991
##
##           Kappa : 0.0278
##           Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.99935
##           Specificity : 0.01567
##           Pos Pred Value : 0.98815
##           Neg Pred Value : 0.22581
##           Prevalence : 0.98797
##           Detection Rate : 0.98733
##           Detection Prevalence : 0.99917
##           Balanced Accuracy : 0.50751
##
##           'Positive' Class : 1
##
```

LDA USING ONLY FULL_TIME_POSITION AND PREVAILING_WAGE

```
lda.fit=lda(CASE_STATUS~.-STATE-occ, data=data.train)
lda.predict = predict(lda.fit, data.test, type = "prob")
confusionMatrix(lda.predict$class, data.test$CASE_STATUS,positive="1")
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction      0      1
##           0      51      1
##           1    6650 550388
##
##           Accuracy : 0.9881
##           95% CI : (0.9878, 0.9883)
##           No Information Rate : 0.988
##           P-Value [Acc > NIR] : 0.2719
##
##           Kappa : 0.0149
##           Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.999998
##           Specificity : 0.007611
##           Pos Pred Value : 0.988062
##           Neg Pred Value : 0.980769
##           Prevalence : 0.987971
##           Detection Rate : 0.987970
##           Detection Prevalence : 0.999907
##           Balanced Accuracy : 0.503804
##
##           'Positive' Class : 1
##
```

LDA USING ONLY FULL_TIME_POSITION, PREVAILING_WAGE AND STATE

```
lda.fit=lda(CASE_STATUS~.-STATE, data=data.train)
lda.predict = predict(lda.fit, data.test, type = "prob")
confusionMatrix(lda.predict$class, data.test$CASE_STATUS,positive="1")
```



```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction      0      1
##           0      99      318
##           1     6602    550071
##
##           Accuracy : 0.9876
##           95% CI : (0.9873, 0.9879)
##           No Information Rate : 0.988
##           P-Value [Acc > NIR] : 0.9964
##
##           Kappa : 0.0264
##           Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.99942
##           Specificity : 0.01477
##           Pos Pred Value : 0.98814
##           Neg Pred Value : 0.23741
##           Prevalence : 0.98797
##           Detection Rate : 0.98740
##           Detection Prevalence : 0.99925
##           Balanced Accuracy : 0.50710
##
##           'Positive' Class : 1
##
```

LDA USING ONLY FULL_TIME_POSITION, PREVAILING_WAGE AND OCCUPATION

```
lda.fit=lda(CASE_STATUS~.-occ, data=data.train)
lda.predict = predict(lda.fit, data.test, type = "prob")
confusionMatrix(lda.predict$class, data.test$CASE_STATUS,positive="1")
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction      0      1
##           0      51      1
##           1    6650 550388
##
##           Accuracy : 0.9881
##           95% CI : (0.9878, 0.9883)
##           No Information Rate : 0.988
##           P-Value [Acc > NIR] : 0.2719
##
##           Kappa : 0.0149
##           Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.999998
##           Specificity : 0.007611
##           Pos Pred Value : 0.988062
##           Neg Pred Value : 0.980769
##           Prevalence : 0.987971
##           Detection Rate : 0.987970
##           Detection Prevalence : 0.999907
##           Balanced Accuracy : 0.503804
##
##           'Positive' Class : 1
##
```

QDA USING USING ALL PREDICTORS

```
qda.fit=qda(CASE_STATUS~., data=data.train)
qda.predict = predict(qda.fit, data.test)
View(qda.predict)
confusionMatrix(qda.predict$class, data.test$CASE_STATUS,positive="1")
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction      0      1
##           0  2428 120754
##           1  4273 429635
##
##           Accuracy : 0.7756
##           95% CI : (0.7745, 0.7767)
##           No Information Rate : 0.988
##           P-Value [Acc > NIR] : 1
##
##           Kappa : 0.0149
##           Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.78060
##           Specificity : 0.36233
##           Pos Pred Value : 0.99015
##           Neg Pred Value : 0.01971
##           Prevalence : 0.98797
##           Detection Rate : 0.77121
##           Detection Prevalence : 0.77888
##           Balanced Accuracy : 0.57147
##
##           'Positive' Class : 1
##
```

QDA USING ONLY FULL_TIME_POSITION AND PREVAILING_WAGE

```
lda.fit=lda(CASE_STATUS~.-STATE-occ, data=data.train)
lda.predict = predict(lda.fit, data.test, type = "prob")
confusionMatrix(lda.predict$class, data.test$CASE_STATUS,positive="1")
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction      0      1
##           0      51      1
##           1    6650 550388
##
##           Accuracy : 0.9881
##           95% CI : (0.9878, 0.9883)
##           No Information Rate : 0.988
##           P-Value [Acc > NIR] : 0.2719
##
##           Kappa : 0.0149
##           Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.999998
##           Specificity : 0.007611
##           Pos Pred Value : 0.988062
##           Neg Pred Value : 0.980769
##           Prevalence : 0.987971
##           Detection Rate : 0.987970
##           Detection Prevalence : 0.999907
##           Balanced Accuracy : 0.503804
##
##           'Positive' Class : 1
##
```

QDA USING ONLY FULL_TIME_POSITION, PREVAILING_WAGE AND STATE

```
lda.fit=lda(CASE_STATUS~.-STATE, data=data.train)
lda.predict = predict(lda.fit, data.test, type = "prob")
confusionMatrix(lda.predict$class, data.test$CASE_STATUS,positive="1")
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction      0      1
##           0      99      318
##           1     6602    550071
##
##           Accuracy : 0.9876
##           95% CI : (0.9873, 0.9879)
##           No Information Rate : 0.988
##           P-Value [Acc > NIR] : 0.9964
##
##           Kappa : 0.0264
##           Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.99942
##           Specificity : 0.01477
##           Pos Pred Value : 0.98814
##           Neg Pred Value : 0.23741
##           Prevalence : 0.98797
##           Detection Rate : 0.98740
##           Detection Prevalence : 0.99925
##           Balanced Accuracy : 0.50710
##
##           'Positive' Class : 1
##
```

QDA USING ONLY FULL_TIME_POSITION, PREVAILING_WAGE AND OCCUPATION

```
lda.fit=lda(CASE_STATUS~.-occ, data=data.train)
lda.predict = predict(lda.fit, data.test, type = "prob")
confusionMatrix(lda.predict$class, data.test$CASE_STATUS,positive="1")
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction      0      1
##           0      51      1
##           1    6650 550388
##
##           Accuracy : 0.9881
##           95% CI : (0.9878, 0.9883)
##           No Information Rate : 0.988
##           P-Value [Acc > NIR] : 0.2719
##
##           Kappa : 0.0149
##           Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.999998
##           Specificity : 0.007611
##           Pos Pred Value : 0.988062
##           Neg Pred Value : 0.980769
##           Prevalence : 0.987971
##           Detection Rate : 0.987970
##           Detection Prevalence : 0.999907
##           Balanced Accuracy : 0.503804
##
##           'Positive' Class : 1
##
```

LOGISTIC REGRESSION USING ALL PREDICTORS and threshold 0.97

```
#Fitting the model on the training dataset
h1bglm.train.fit =glm(CASE_STATUS~., family=binomial(link = logit), data = data.train
)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```


```
#Finding Predicitons on Testing set
prediction=predict(h1bglm.train.fit, newdata=data.test, type="response")

#Threshold
pred_class = vector()
pred_class[prediction<0.97]=0
pred_class[prediction>=0.97]=1

##confusion matrix
confusionMatrix(pred_class, data.test$CASE_STATUS,positive="1")
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction      0      1
##           0  2213  73858
##           1  4488 476531
##
##           Accuracy : 0.8594
##           95% CI : (0.8584, 0.8603)
##           No Information Rate : 0.988
##           P-Value [Acc > NIR] : 1
##
##           Kappa : 0.0321
##           Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.86581
##           Specificity : 0.33025
##           Pos Pred Value : 0.99067
##           Neg Pred Value : 0.02909
##           Prevalence : 0.98797
##           Detection Rate : 0.85539
##           Detection Prevalence : 0.86345
##           Balanced Accuracy : 0.59803
##
##           'Positive' Class : 1
##
```

```
#ROC Curve
pred = prediction(pred_class, data.test$CASE_STATUS)
perf = performance(pred, "tpr", "fpr")
plot(perf)
```



```
#Area Under the Curve  
auc.tmp = performance(pred,"auc");  
auc = as.numeric(auc.tmp@y.values)  
auc
```

```
## [1] 0.5980284
```

LOGISTIC REGRESSION USING ONLY FULL_TIME_POSITION AND PREVAILING_WAGE

```
#Fitting the model on the training dataset  
h1bglm.train.fit =glm(CASE_STATUS~.-STATE-occ, family=binomial(link = logit), data =  
data.train)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```



```

#Finding Prdicitons on Testing set
prediction=predict(h1bglm.train.fit, newdata=data.test, type="response")

#Threshold
pred_class = vector()
pred_class[prediction<0.97]=0
pred_class[prediction>=0.97]=1

##confusion matrix
confusionMatrix(pred_class, data.test$CASE_STATUS,positive="1")

```

```


## Confusion Matrix and Statistics
##
##           Reference
## Prediction      0      1
##           0  3566 289791
##           1  3135 260598
##
##           Accuracy : 0.4742
##           95% CI : (0.4729, 0.4755)
##           No Information Rate : 0.988
##           P-Value [Acc > NIR] : 1
##
##           Kappa : 3e-04
##           McNemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.47348
##           Specificity : 0.53216
##           Pos Pred Value : 0.98811
##           Neg Pred Value : 0.01216
##           Prevalence : 0.98797
##           Detection Rate : 0.46778
##           Detection Prevalence : 0.47341
##           Balanced Accuracy : 0.50282
##
##           'Positive' Class : 1
##

```

```

#ROC Curve
pred = prediction(pred_class, data.test$CASE_STATUS)
perf = performance(pred,"tpr","fpr")
plot(perf)

```



```
#Area Under the Curve  
auc.tmp = performance(pred,"auc");  
auc = as.numeric(auc.tmp@y.values)  
auc
```

```
## [1] 0.5028195
```

LOGISTIC REGRESSION USING ONLY FULL_TIME_POSITION, PREVAILING_WAGE AND STATE

```
#Fitting the model on the training dataset  
h1bglm.train.fit =glm(CASE_STATUS~.-occ, family=binomial(link = logit), data = data.t  
rain)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```

#Finding Prdicitons on Testing set
prediction=predict(h1bglm.train.fit, newdata=data.test, type="response")

#Threshold
pred_class = vector()
pred_class[prediction<0.97]=0
pred_class[prediction>=0.97]=1

##confusion matrix
confusionMatrix(pred_class, data.test$CASE_STATUS,positive="1")

```

```


## Confusion Matrix and Statistics
##
##           Reference
## Prediction      0      1
##           0  3512 281733
##           1  3189 268656
##
##           Accuracy : 0.4886
##           95% CI : (0.4872, 0.4899)
##       No Information Rate : 0.988
##       P-Value [Acc > NIR] : 1
##
##           Kappa : 6e-04
##  McNemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.48812
##           Specificity : 0.52410
##           Pos Pred Value : 0.98827
##           Neg Pred Value : 0.01231
##           Prevalence : 0.98797
##           Detection Rate : 0.48225
##       Detection Prevalence : 0.48797
##           Balanced Accuracy : 0.50611
##
##           'Positive' Class : 1
##

```

```

#ROC Curve
pred = prediction(pred_class, data.test$CASE_STATUS)
perf = performance(pred,"tpr","fpr")
plot(perf)

```



```
#Area Under the Curve
auc.tmp = performance(pred,"auc");
auc = as.numeric(auc.tmp@y.values)
auc
```

```
## [1] 0.5061106
```

LOGISTIC REGRESSION USING ONLY FULL_TIME_POSITION, PREVAILING_WAGE AND OCCUPATION

```
#Fitting the model on the training dataset
h1bglm.train.fit =glm(CASE_STATUS~.-STATE, family=binomial(link = logit), data = dat
a.train)
```

```
## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```

#Finding Prdicitons on Testing set
prediction=predict(h1bglm.train.fit, newdata=data.test, type="response")

#Threshold
pred_class = vector()
pred_class[prediction<0.97]=0
pred_class[prediction>=0.97]=1

#confusion matrix
confusionMatrix(pred_class, data.test$CASE_STATUS,positive="1")

```

```


## Confusion Matrix and Statistics
##
##           Reference
## Prediction      0      1
##           0  2246  78029
##           1  4455  472360
##
##           Accuracy : 0.8519
##           95% CI : (0.851, 0.8529)
##       No Information Rate : 0.988
##       P-Value [Acc > NIR] : 1
##
##           Kappa : 0.0301
##  McNemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.85823
##           Specificity : 0.33517
##           Pos Pred Value : 0.99066
##           Neg Pred Value : 0.02798
##           Prevalence : 0.98797
##           Detection Rate : 0.84791
##       Detection Prevalence : 0.85590
##           Balanced Accuracy : 0.59670
##
##           'Positive' Class : 1
##

```

```

#ROC Curve
pred = prediction(pred_class, data.test$CASE_STATUS)
perf = performance(pred,"tpr","fpr")
plot(perf)

```



```
#Area Under the Curve
auc.tmp = performance(pred,"auc");
auc = as.numeric(auc.tmp@y.values)
auc
```

```
## [1] 0.5967016
```

DECISION TREE USING ALL PREDICTORS AND THRESHOLD 0.97


```
ctree1=ctree(CASE_STATUS~., data=data.train)
probabilities = 1-unlist(treeresponse(ctree1,newdata=data.test), use.names=F)[seq(1,n
row(data.test)*2,2)]

probabilities[probabilities<0.97]=0
probabilities[probabilities>=0.97]=1

confusionMatrix(probabilities, data.test$CASE_STATUS,positive="1")
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction      0      1
##           0  2194  87316
##           1  4507  463073
##
##           Accuracy : 0.8352
##           95% CI : (0.8342, 0.8361)
##           No Information Rate : 0.988
##           P-Value [Acc > NIR] : 1
##
##           Kappa : 0.0238
##           Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.84136
##           Specificity : 0.32741
##           Pos Pred Value : 0.99036
##           Neg Pred Value : 0.02451
##           Prevalence : 0.98797
##           Detection Rate : 0.83124
##           Detection Prevalence : 0.83933
##           Balanced Accuracy : 0.58438
##
##           'Positive' Class : 1
##
```

```
#ROC CURVE
pred1 = prediction( probabilities, data.test$CASE_STATUS)
perf1 = performance(pred1,"tpr","fpr")
plot(perf1)
```



```
#AREA Under the Curve
auc.tmp1 = performance(pred1,"auc");
auc1 = as.numeric(auc.tmp1@y.values)
auc1
```

```
## [1] 0.5843848
```

DECISION TREE USING ALL PREDICTORS AND THRESHOLD 0.5

```
ctree1=ctree(CASE_STATUS~., data=data.train)
probabilities = 1-unlist(treeresponse(ctree1,newdata=data.test), use.names=F)[seq(1,n
row(data.test)*2,2)]


probabilities[probabilities<0.5]=0
probabilities[probabilities>=0.5]=1

confusionMatrix(probabilities, data.test$CASE_STATUS,positive="1")
```



```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction      0      1
##           0      79      15
##           1    6622  550374
##
##           Accuracy : 0.9881
##           95% CI : (0.9878, 0.9884)
##           No Information Rate : 0.988
##           P-Value [Acc > NIR] : 0.2178
##
##           Kappa : 0.0229
##           Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.99997
##           Specificity : 0.01179
##           Pos Pred Value : 0.98811
##           Neg Pred Value : 0.84043
##           Prevalence : 0.98797
##           Detection Rate : 0.98794
##           Detection Prevalence : 0.99983
##           Balanced Accuracy : 0.50588
##
##           'Positive' Class : 1
##
```

```
#ROC CURVE
pred1 = prediction( probabilities, data.test$CASE_STATUS)
perf1 = performance(pred1,"tpr","fpr")
plot(perf1)
```



```
#AREA Under the Curve
auc.tmp1 = performance(pred1,"auc");
auc1 = as.numeric(auc.tmp1@y.values)
auc1
```

```
## [1] 0.505881
```

DECISION TREE USING ONLY FULL_TIME_POSITION AND PREVAILING_WAGE


```
ctree1=ctree(CASE_STATUS~.-STATE-occ, data=data.train)
probabilities = 1-unlist(treeresponse(ctree1,newdata=data.test), use.names=F)[seq(1,n
row(data.test)*2,2)]

probabilities[probabilities<0.97]=0
probabilities[probabilities>=0.97]=1

confusionMatrix(probabilities, data.test$CASE_STATUS,positive="1")
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction      0      1
##           0  2194  87316
##           1  4507  463073
##
##           Accuracy : 0.8352
##           95% CI : (0.8342, 0.8361)
##           No Information Rate : 0.988
##           P-Value [Acc > NIR] : 1
##
##           Kappa : 0.0238
##           Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.84136
##           Specificity : 0.32741
##           Pos Pred Value : 0.99036
##           Neg Pred Value : 0.02451
##           Prevalence : 0.98797
##           Detection Rate : 0.83124
##           Detection Prevalence : 0.83933
##           Balanced Accuracy : 0.58438
##
##           'Positive' Class : 1
##
```

```
#ROC CURVE
pred1 = prediction( probabilities, data.test$CASE_STATUS)
perf1 = performance(pred1,"tpr","fpr")
plot(perf1)
```



```
#AREA Under the Curve
auc.tmp1 = performance(pred1,"auc");
auc1 = as.numeric(auc.tmp1@y.values)
auc1
```

```
## [1] 0.5843848
```

DECISION TREE USING ONLY FULL_TIME_POSITION, PREVAILING_WAGE AND STATE


```
ctree1=ctree(CASE_STATUS~.-occ, data=data.train)
probabilities = 1-unlist(treeresponse(ctree1,newdata=data.test), use.names=F)[seq(1,n
row(data.test)*2,2)]

probabilities[probabilities<0.97]=0
probabilities[probabilities>=0.97]=1

confusionMatrix(probabilities, data.test$CASE_STATUS,positive="1")
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction      0      1
##           0  2194  87316
##           1  4507  463073
##
##           Accuracy : 0.8352
##           95% CI : (0.8342, 0.8361)
##           No Information Rate : 0.988
##           P-Value [Acc > NIR] : 1
##
##           Kappa : 0.0238
##           Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.84136
##           Specificity : 0.32741
##           Pos Pred Value : 0.99036
##           Neg Pred Value : 0.02451
##           Prevalence : 0.98797
##           Detection Rate : 0.83124
##           Detection Prevalence : 0.83933
##           Balanced Accuracy : 0.58438
##
##           'Positive' Class : 1
##
```

```
#ROC CURVE
pred1 = prediction( probabilities, data.test$CASE_STATUS)
perf1 = performance(pred1,"tpr","fpr")
plot(perf1)
```



```
#AREA Under the Curve
auc.tmp1 = performance(pred1,"auc");
auc1 = as.numeric(auc.tmp1@y.values)
auc1
```

```
## [1] 0.5843848
```

DECISION TREE USING ONLY FULL_TIME_POSITION, PREVAILING_WAGE AND OCCUPATION


```
ctree1=ctree(CASE_STATUS~.-STATE, data=data.train)
probabilities = 1-unlist(treeresponse(ctree1,newdata=data.test), use.names=F)[seq(1,n
row(data.test)*2,2)]

probabilities[probabilities<0.97]=0
probabilities[probabilities>=0.97]=1

confusionMatrix(probabilities, data.test$CASE_STATUS,positive="1")
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction      0      1
##           0  2194  87316
##           1  4507  463073
##
##           Accuracy : 0.8352
##           95% CI : (0.8342, 0.8361)
##           No Information Rate : 0.988
##           P-Value [Acc > NIR] : 1
##
##           Kappa : 0.0238
##           Mcnemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.84136
##           Specificity : 0.32741
##           Pos Pred Value : 0.99036
##           Neg Pred Value : 0.02451
##           Prevalence : 0.98797
##           Detection Rate : 0.83124
##           Detection Prevalence : 0.83933
##           Balanced Accuracy : 0.58438
##
##           'Positive' Class : 1
##
```

```
#ROC CURVE
pred1 = prediction( probabilities, data.test$CASE_STATUS)
perf1 = performance(pred1,"tpr","fpr")
plot(perf1)
```



```
#AREA Under the Curve  
auc.tmp1 = performance(pred1,"auc");  
auc1 = as.numeric(auc.tmp1@y.values)  
auc1
```

```
## [1] 0.5843848
```